

Networking

Greg Walker 16 Jun 2024 Download PDF

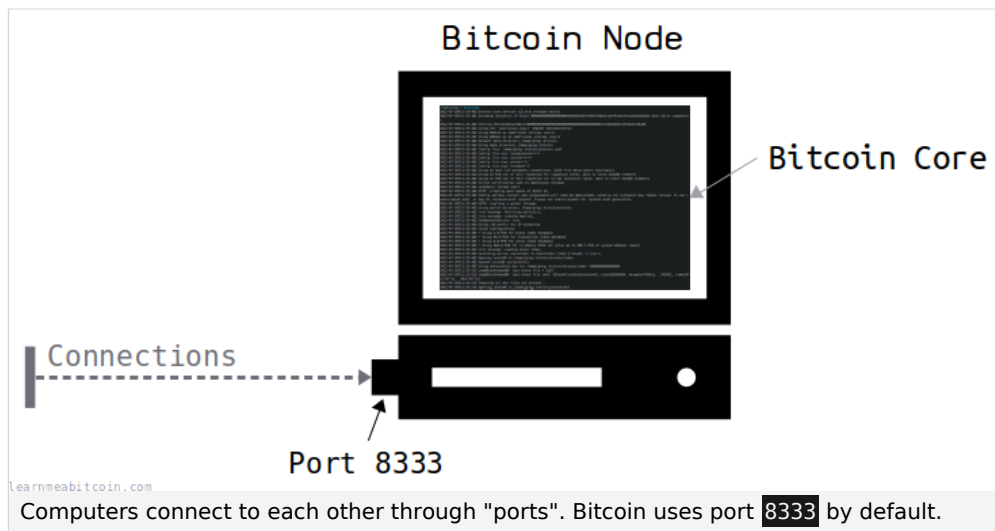
Here's a quick guide on how to *connect to* and *communicate with* a node on the Bitcoin network.



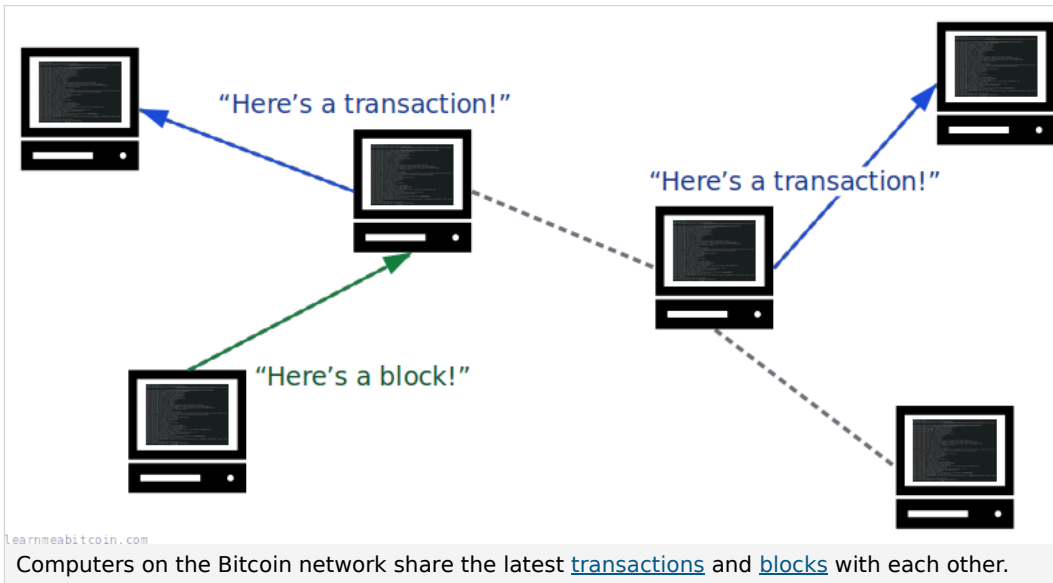
0. Intro

Bitcoin is a computer program. You can [download](#) it for free.

It runs on an open port on your computer, which means anyone can to connect it and communicate with it across the Internet.



When you run Bitcoin, it uses ports to connect to other computers running the same program. So when you have lots of people running Bitcoin, you end up with a network of computers connected together and communicating with each other.

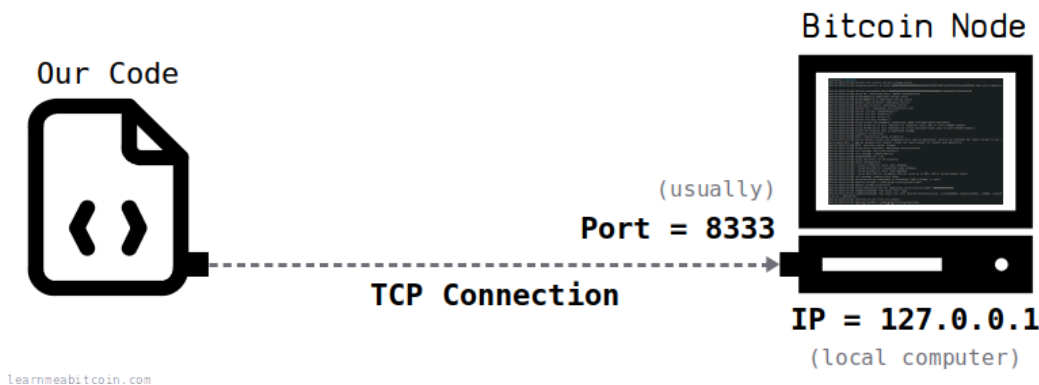


Anyway, the cool thing about Bitcoin is **you can write your own basic program to connect to a node if you want to**. You just need to know how to speak its language.

In this guide I'm going to show you how to connect to a bitcoin node using [Ruby 2](#). Ruby is a simple language, so you should be able to translate the code in to whichever language you prefer to use. I personally like Ruby.

And trust me, if I can connect to a Bitcoin node, anyone can.

1. Connecting



First things first, two quick facts you need to know about the Bitcoin program:

- It runs on **port 8333** (usually)
- It uses **TCP** for communication

So all you need to connect to a Bitcoin node is the **IP address** of the computer it's running on, and the ability to make **TCP connections** from your programming language. For example:

```
Ruby

# Sockets are in the standard library in Ruby
require 'socket'

# Open a TCP connection to an IP and port
socket = TCPSocket.open("162.120.69.182", 8333) # local computer = 127.0.0.1
```

But that's pretty boring on its own. To start *receiving* data (like actual [transactions](#) and [blocks](#)), you need to start by sending it some *messages* first.

- See [Finding Nodes](#) if you don't already have an IP to connect to. The easiest method is to connect to your own local node (`127.0.0.1`), or you could try connecting to the node running on this server if you prefer (`162.120.69.182`).
- You can use this [Bitnodes.io tool](#) to check if a remote node is accepting incoming connections.

```
mainnet = 8333
testnet = 18333
regtest = 18444
```

2. Messages

The diagram illustrates the structure of a message and its transmission over a TCP connection. On the left, a document icon with code symbols (<>) represents the message. The message is structured as follows:

message	
header	
payload	

A dashed arrow labeled "TCP Connection" points from the message structure to a computer icon on the right, which displays a terminal window with code, representing the destination system.

[illegible]

So I'm going to start by showing you the basic structure of a message `header` and `payload`, and then I'll show you how to construct one yourself. I'm going to use a "version" type message as the first example, as that's the first message you want to send to a Bitcoin node after connecting to one.

The “version” message provides information about the transmitting node to the receiving node at the beginning of a connection. Until both peers have exchanged “version” messages, no other messages will be accepted.

– developer.bitcoin.org

Version

Header

The header contains a **summary of the message**, and its structure is the same for every message in the Bitcoin protocol.

Here's what a header looks like for a "version" message:

Header: (version message)

Name	Example Data	Format	Size	Bytes
Magic Bytes		bytes	4	F9 BE B4 D9
Command	"version"	ascii bytes	12	76 65 72 73 69 6F 6E 00 00 00 00 00
Size	85	little-endian	4	55 00 00 00
Checksum		bytes	4	F7 63 9C 60

Fields

- **Magic Bytes:** This is a unique set of bytes used to identify the start of a new message. They're always the same. You see, you'll be reading a stream of bytes from your TCP connection when receiving messages, so it's handy to be able to identify when a new message starts. This random-looking set of bytes has been specifically chosen so that it's unlikely that they would appear somewhere else in a message.
- **Command:** This indicates the type of message being sent. You can send different types of messages in the Bitcoin protocol, and they contain different types of information. It's a 12-byte field containing the [ASCII \[+ \]](#) encoding of the name of the message type. The one in this example says that we are sending a "version" message, which is used to send information about ourselves to another node.
- **Size:** This is the size of the upcoming payload. This indicates how many bytes you need to read from the socket to get the full message being sent.
- **Checksum:** This is a small fingerprint for the payload. It allows us to quickly check that the data in the payload hasn't been tampered with during transit. It's created by double-hashing the payload, then taking the first 4 bytes of the result.

Payload

The payload contains the main **content of the message**. Different message types have different structures for their payloads.

Here's the payload for a "version" message:

Payload (version message):

Name	Example Data	Format	Size	Example Bytes
Protocol Version	70014	little-endian	4	7E 11 01 00
Services	0	bit field , little-endian	8	00 00 00 00 00 00 00 00
Time	1640961477	little-endian	8	C5 15 CF 61 00 00 00 00
Remote Services	0	bit field, little-endian	8	00 00 00 00 00 00 00 00
Remote IP	46.19.137.74	ipv6, big-endian	16	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
Remote Port	8333	big-endian	2	20 8D
Local Services	0	bit field, little-endian	8	00 00 00 00 00 00 00 00
local IP	127.0.0.1	ipv6, big-endian	16	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
Local Port	8333	big-endian	2	20 8D

Nonce	0	little-endian	8	00 00 00 00 00 00 00 00
User Agent	" "	compact size , ascii	compact	00
Last Block	0	little-endian	4	00 00 00 00

A "version" message is one of the more complex messages you can send in Bitcoin, but that's only because it contains lots of information. It's a good place to start though, because if you can construct a "version" message, you can construct any message in the Bitcoin protocol.

Fields

Here's what each of the individual fields mean for this particular message:

- **Protocol Version:** This is the version of the protocol our node understands. Different versions of the protocol have different messages, so by giving our protocol version we let the other node know what kind of messages we can work with.
- **Services:** This is a list of optional services that your node can offer. This is a 32 bit field, where each bit can be set to **1** to indicate a different service you offer. For example, setting the first bit (on the right) indicates that you are a full node and can provide all of the blocks in the blockchain. You can leave this as zero if you're just testing.
- **Time:** Your computer's time as a Unix timestamp (the number of seconds since 01 Jan 1970).
- **Remote Services:** This is the list of optional services you think the node you're connecting to can offer. It's the same structure as the main "Services" field above. I'm not sure why this is useful or if it's actually used for anything, so I leave it as zero.
- **Remote IP:** This is the IP address of the node you think you're connecting to. This is in IPv6 format (you can easily [convert between IPv4 and IPv6](#) if you need to). I don't think this is crucial either, but I set it as the IP I'm connecting to anyway.
- **Remote Port:** This is the port on the node you think you're connecting to. I just leave it as the default **8333**.
- **Local Services:** This is the list of services your node offers. I'm not sure why this gets repeated.
- **Local IP:** This is what you think your local IP is. This is your IP address in IPv6 format (again, you can [convert between IPv4 and IPv6](#) if you need to). This is not actually used by the remote node, so you can set it to what you want. I just set it as localhost (**127.0.0.1**).
- **Local Port:** This is the local port you're communicating from. Again, I don't think this is crucial, but I leave it as the default **8333**.
- **Nonce:** A randomly generated number that be used to detect connections to yourself later on. You can leave this as zero if it's not needed and it will just be ignored. The term "nonce" is short for "number used once", unless you're British, in which case it doesn't mean that at all.
- **User Agent.** A custom string you can use to identify the make and model of your node on the network. Bitcoin Core uses a string like "/Satoshi:22.0.0/", but you can put "Awesome Node 5000" or something like that if you prefer. You can see these user agents for yourself when you run `bitcoin-cli getpeerinfo`. You can leave this field blank if you want, but just remember you'll still need to place a **00** byte in this field to indicate that you have not provided any upcoming bytes.
- **Last Block:** The height of the top block in your local blockchain. Leave this as zero if you haven't got any blocks or don't care to share any.

As I say, this is one of the more complicated messages, so don't let it put you off from trying to connect to a node. Give it a go.

Code [hide]

Here's some example code for constructing a "version" message in Ruby:

 Ruby


```
require 'digest' # needed for creating checksums

# Handy functions for getting data in the right format for messages
def hexadecimal(number)
  return number.to_s(16)
```


So now we now how to construct a message, we can start communicating with the node we've just connected to.

Handshaking is the process that establishes communication between two networking devices.

In the Bitcoin protocol, the handshake works like this:



- And that's all there is to it.

We need to send two messages to perform the handshake:

- We've already prepared our "version" message, so let's create a "verack" message.

The "verack" is a simple message header without a payload:

<https://learnmeabitcoin.com/technical/networking/>

Name	Example Data	Format	Size	Example Bytes
Magic Bytes		bytes	4	F9 BE B4 D9
Command	"verack"	ascii bytes	12	76 65 72 61 63 6B 00 00 00 00 00 00
Size	0	little-endian	0	00 00 00 00
Checksum		bytes	4	5D F6 E0 E2

Hexadecimal: F9BEB4D976657261636B000000000000000000005DF6E0E2

Tip: A "verack" message is always the same.

Sending and receiving messages

Now we've got our messages ready, we just need to send them to the node we've connected to (and receive messages back from them).

- To "send" messages, we just *write* bytes to our TCP socket connection.
- To "receive" messages we just *read* bytes from the same socket.

Code [hide]

Here's some Ruby code showing how to manually construct each message, and how to write/read bytes to/from the socket connection:

Ruby

```
# 1. Send Version Message

# Prepare version message
version = message

# Write the message to the socket (the protocol sends and receives messages in raw bytes)
socket.write [version].pack("H*")
puts "version->"
puts version
puts

# 2. Receive Version Message

# Read the message header response from the socket
magic_bytes = socket.read(4)
command      = socket.read(12)
size         = socket.read(4)
checksum     = socket.read(4)

# View the message header
puts "<-version"
puts "magic_bytes: " + magic_bytes.unpack("H*").join # convert raw bytes to hexadecimal characters
puts "command:     " + command.to_s                 # to_s automatically converts raw bytes to ASCII characters
puts "size:        " + size.unpack("V").join         # V = 32-byte unsigned, little-endian
puts "checksum:    " + checksum.unpack("H*").join

# Read the message payload
size = size.unpack("V").join.to_i
payload = socket.read(size)
```



```

# View the message payload
puts "payload: " + payload.unpack("H*").join
puts

# 3. Receive Verack Message (verack = version acknowledged)

# Read the message header response from the socket
magic_bytes = socket.read(4)
command     = socket.read(12)
size        = socket.read(4)
checksum    = socket.read(4)

# View the message header
puts "<-verack"
puts "magic_bytes: " + magic_bytes.unpack("H*").join # convert raw bytes to hexadecimal characters
puts "command: " + command.to_s # to_s automatically converts raw bytes to ASCII characters
puts "size: " + size.unpack("V").join # V = 32-byte unsigned, little-endian
puts "checksum: " + checksum.unpack("H*").join

# Read the message payload (there shouldn't be any)
size = size.unpack("V").join.to_i
payload = socket.read(size)

# View the message payload (there shouldn't be any)
puts "payload: " + payload.unpack("H*").join
puts

# 4. Send Verack Message

# Create verack message
payload = '' # verack has no payload, it's just a message header
magic_bytes = 'f9beb4d9'
command = ascii2hex('verack')
size = reversebytes(size(hexadecimal(payload.size/2), 4))
checksum = checksum(payload)
verack = magic_bytes + command + size + checksum + payload

# Write the message to the socket
socket.write [verack].pack("H*")
puts "verack->"
puts "magic_bytes: " + magic_bytes
puts "command: " + 'verack'
puts "size: " + size.to_i(16).to_s
puts "checksum: " + checksum
puts "payload: " + payload
puts

```

Note: Strings and Bytes. When sending data "over the wire" you need to convert all of your data to raw bytes. In the code examples I've given, even though it looks like I'm working with bytes, I'm actually manipulating *strings* made up of hexadecimal characters that *represent* bytes. That's where the `pack()` function comes in, as this allows you to convert strings to actual bytes. Your programming language will have something similar.


Note: Socket Programming. The way you write/read bytes to/from a socket will be different from one programming language to another, so it might take some getting used to if you've never done it before.

Anyway, once you've received that "verack" message (and sent your own one back), the handshake is complete. And if everything has worked correctly, the node will start sending you some *new* message types...

4. Receiving Messages

The node we've just connected to will continuously send us new messages after the handshake. So to keep receiving these messages, all we need to do is **keep reading from the socket in a loop**.

This is what the new messages are going to look like:



learnmeabitcoin.com

Note: You may receive some different messages before the "inv" messages shown in the diagram above, depending on what protocol version you're using. I'm just going to ignore them for now as they're not critically important.

I'll explain what these "inv" messages are and how to respond to them in a moment. But for now I'll just show you how to *keep reading messages* from the node you've connected to:

Code [hide]

The following code is similar to the code from before, except this time we've put it in a *loop* to continuously read from the socket.

 Ruby

```

# Keep reading messages
loop do

  # Create an empty buffer to help us find the next stream of magic bytes (the start of a new message)
  buffer = ''

  # Keep looping to read bytes from the socket
  loop do

    # Read one byte at the time
    byte = socket.read(1)

    # Check that we haven't been disconnected from the node.
    if byte.nil?
      puts "Read a nil byte from the socket. Looks like the remote node has disconnected from us. We prob
      exit
    end

    # Add each byte to the temporary buffer
  end
end
  
```

```

buffer += byte.unpack("H*").join unless byte.nil? # do not do anything if we got a nil byte for some reason

# Check the buffer when it reaches 4 bytes
if (buffer.size == 8) # 8 hexadecimal characters = 4 bytes

  # See if the buffer matches the magic bytes
  if (buffer == 'f9beb4d9')

    # If we've got the magic bytes we're looking for, go ahead and read the full message from the socket
    command = socket.read(12).to_s.delete("\x00") # convert to ascii and remove any empty bytes
    size = socket.read(4).unpack("V").join.to_i # convert to an integer
    checksum = socket.read(4).unpack("H*").join # convert to hexadecimal string of bytes
    payload = socket.read(size).unpack("H*").join # use the size from the header to read the payload

    # Print the message
    puts "<-#{command}"
    puts "magic_bytes: " + buffer
    puts "command: " + command
    puts "size: " + size.to_s
    puts "checksum: " + checksum
    puts "payload: " + payload
    puts

    # Break out of the loop for reading a single message
    break

  end

  # Reset the buffer and keep looking for a stream of magic bytes
  buffer = ''

end

end

end
end

```


Now we can keep reading data from this node forever, or at least until our computer randomly crashes and I end up losing all the code I've been writing for the last hour because I forgot to save it.

5. Requesting Transactions and Blocks

A node won't openly send you all the new transactions and blocks that it has received. Instead, to save bandwidth, they will send you a *list* of hashes of the latest transactions and blocks they've received in "inv" (inventory) messages.

You can then respond to these "inv" messages listing all the specific transactions and blocks you want with "getdata" messages.

Then, after you've sent your "getdata" message, the node will send you the full transactions and blocks you've requested in subsequent "tx" and "block" messages:



Inv

The payload of an "inv" message looks like this:

Payload: (inv)

Name	Format	Size	Example Bytes
Count	compact size	variable	01
Inventory	inventory vector	variable	01 00 00 00 aa 32 5e 91 22 aa 39 ca 18 c7 5a ab e2 a3 ce af 98 02

Inventory

The "Inventory" part of the payload is *another* data structure in itself. But it's pretty simple: it's just a list of [transaction hashes](#) and/or [block hashes](#):

Inventory:

Name	Format	Size	Example Bytes
Type	little-endian	4	01 00 00 00
Hash	bytes	32	aa 32 5e 91 22 aa 39 ca 18 c7 5a ab e2 a3 ce af 98 02 ac d1 a4 07 20 92 5b

Types:

- * 01 00 00 00 = MSG_TX (Transaction Hash)
- * 02 00 00 00 = MSG_BLOCK (Block Hash)

Getdata

The "getdata" message you respond with has the exact same structure as the "inv" message (which is convenient):

Payload: (getdata)

Name	Format	Size	Example Bytes
Count	compact size	variable	01

Inventory	inventory vector	variable	01 00 00 00 aa 32 5e 91 22 aa 39 ca 18 c7 5a ab e2 a3 ce af 98 02
-----------	------------------	----------	---

So if you want *all* of the transactions and blocks in the "inv", you can just reply with the exact same payload in your "getdata" message. Or if you don't want them all, just construct a payload with a list of the transaction/block hashes that you do want.

Note:

SegWit Transactions

To request the full transaction data for new [segwit transactions](#) (i.e. including the [witness](#) data), you must change the *type* field in the [inventory](#) part of your "getdata" message from:

- 01 00 00 00 = MSG_TX
- 02 00 00 00 = MSG_BLOCK

To:

- 01 00 00 40 = MSG_WITNESS_TX
- 02 00 00 40 = MSG_WITNESS_BLOCK

For example, the payload for the example "getdata" message above would be:

Payload: (getdata)			
Name	Format	Size	Example Bytes
Count	compact size	variable	01
Inventory	inventory vector	variable	01 00 00 40 aa 32 5e 91 22 aa 39 ca 18 c7 5a ab e2 a3 ce af 98 02 ac d1 a4 07 20 92 5b fd 77 ff f

You should make this change for all "getdata" messages to make sure you're getting the full transaction data for both segwit and legacy transactions.

Anyway, after sending your "getdata" message, the node will proceed to send you full copies of the transactions and blocks you asked for in individual "tx" and "block" messages in response.

Code [hide]

Here's some Ruby code that responds to every "inv" with a "getdata" message requesting everything in the payload:

 Ruby

```
# Keep reading messages
loop do

  # Create an empty buffer to help us find the next stream of magic bytes (the start of a new message)
  buffer = ''

  # Keep looping to read bytes from the socket
  loop do

    # Read one byte at the time
    byte = socket.read(1)

    # Check that we haven't been disconnected from the node.
    if byte.nil?
      puts "Read a nil byte from the socket. Looks like the remote node has disconnected from us. We pro
      exit
    end
```

```

# Add each byte to the temporary buffer
buffer += byte.unpack("H*").join unless byte.nil? # do not do anything if we got a nil byte for some reason

# Check the buffer when it reaches 4 bytes
if (buffer.size == 8) # 8 hexadecimal characters = 4 bytes

  # See if the buffer matches the magic bytes
  if (buffer == 'f9beb4d9')

    # If we've got the magic bytes we're looking for, go ahead and read the full message from the socket
    command = socket.read(12).to_s.delete("\x00") # convert to ascii and remove any empty bytes
    size = socket.read(4).unpack("V").join.to_i # convert to an integer
    checksum = socket.read(4).unpack("H*").join # convert to hexadecimal string of bytes
    payload = socket.read(size).unpack("H*").join # use the size from the header to read the payload

    # Print the message
    puts "<#{command}"
    puts "magic_bytes: " + buffer
    puts "command: " + command
    puts "size: " + size.to_s
    puts "checksum: " + checksum
    puts "payload: " + payload
    puts

    # Respond to all inv messages with getdata messages
    if command == "inv"

      # Set new command name
      command = "getdata"

      # Use the same payload as the one we got from the inv message
      payload = payload

      # Create message
      magic_bytes = 'f9beb4d9'
      command_hex = ascii2hex(command)
      size = reversebytes(size(hexadecimal(payload.size/2), 4))
      checksum = checksum(payload)
      message = magic_bytes + command_hex + size + checksum + payload

      # Print the message header and payload
      puts "#{command}->"
      puts "magic_bytes: " + magic_bytes
      puts "command: " + command
      puts "size: " + (payload.size/2).to_s
      puts "checksum: " + checksum
      puts "payload: " + payload
      puts

      # Send the message (convert from hexadecimal string to raw bytes first)
      socket.write [message].pack("H*")

    end

    # Break out of the loop for reading a single message
    break

  end

end

# Reset the buffer and keep looking for a stream of magic bytes

```

```

    buffer = ''

    end

    end

    end

```


And that's how you can get the latest transactions and blocks from an actual node on the network.

If you've got this far and everything is working, you've figured out how to connect to and communicate with a bitcoin node from scratch. Everything from here just involves constructing different types of messages.

Note: Here's a [full list](#) of the messages Bitcoin nodes can send each other.

6. Keeping Connected

One last thing before you go: the node you've just connected to will occasionally send you "ping" messages to see if you're still there. So if you want to keep the connection alive, you'll need to respond with timely "pong" messages.



learnmeabitcoin.com

Ping

As of protocol version **60001**, each "ping" message contains a random number as its payload:

Payload: (ping)

Name	Format	Size	Example Bytes
Nonce	bytes	8	88 c8 49 39 65 b6 41 69

Pong


Your "pong" message in response just needs to contain with the same number in its payload too:

Payload: (pong)

Name	Format	Size	Example Bytes
Nonce	bytes	8	88 c8 49 39 65 b6 41 69

So by adding one last adjustment to our loop, we can now keep the connection open and receive transactions and blocks forever:

Code [hide]

 Ruby

```
# Keep reading messages
loop do

  # Create an empty buffer to help us find the next stream of magic bytes (the start of a new message)
  buffer = ''

  # Keep looping to read bytes from the socket
  loop do

    # Read one byte at the time
    byte = socket.read(1)

    # Check that we haven't been disconnected from the node.
    if byte.nil?
      puts "Read a nil byte from the socket. Looks like the remote node has disconnected from us. We prob
      exit
    end

    # Add each byte to the temporary buffer
    buffer += byte.unpack("H*").join unless byte.nil? # do not do anything if we got a nil byte for some r

    # Check the buffer when it reaches 4 bytes
    if (buffer.size == 8) # 8 hexadecimal characters = 4 bytes

      # See if the buffer matches the magic bytes
      if (buffer == 'f9beb4d9')

        # If we've got the magic bytes we're looking for, go ahead and read the full message from the s
        command      = socket.read(12).to_s.delete("\x00") # convert to ascii and remove any empty byte
        size         = socket.read(4).unpack("V").join.to_i # convert to an integer
        checksum     = socket.read(4).unpack("H*").join    # convert to hexadecimal string of bytes
        payload      = socket.read(size).unpack("H*").join  # use the size from the header to read the p

        # Print the message
        puts "<#{command}"
        puts "magic_bytes: " + buffer
        puts "command:      " + command
        puts "size:         " + size.to_s
        puts "checksum:     " + checksum
        puts "payload:      " + payload
        puts

        # Respond to all inv messages with getdata messages
        if command == "inv"

          # Set new command name
          command = "getdata"

          # Use the same payload as the one we got from the inv message
          payload = payload

          # Create message
          magic_bytes = 'f9beb4d9'
```



```

command_hex = ascii2hex(command)
size         = reversebytes(size(hexadecimal(payload.size/2), 4))
checksum     = checksum(payload)
message      = magic_bytes + command_hex + size + checksum + payload

# Print the message header and payload
puts "#{command}->"
puts "magic_bytes: " + magic_bytes
puts "command:      " + command
puts "size:         " + (payload.size/2).to_s
puts "checksum:     " + checksum
puts "payload:      " + payload
puts

# Send the message (convert from hexadecimal string to raw bytes first)
socket.write [message].pack("H*")

end

# Respond to all ping messages with pong messages
if command == "ping"

  # Set new command name
  command = "pong"

  # Use the same payload as the one we got from the ping message
  payload = payload

  # Create message
  magic_bytes = 'f9beb4d9'
  command_hex = ascii2hex(command)
  size         = reversebytes(size(hexadecimal(payload.size/2), 4))
  checksum     = checksum(payload)
  message      = magic_bytes + command_hex + size + checksum + payload

  # Print the message header and payload
  puts "#{command}->"
  puts "magic_bytes: " + magic_bytes
  puts "command:      " + command
  puts "size:         " + (payload.size/2).to_s
  puts "checksum:     " + checksum
  puts "payload:      " + payload
  puts

  # Send the message (convert from hexadecimal string to raw bytes first)
  socket.write [message].pack("H*")

end

# Break out of the loop for reading a single message
break

end

# Reset the buffer and keep looking for a stream of magic bytes
buffer = ''

end

end

```

```
end
```

Note: Functions. I've repeated the same code in my code examples to keep everything as readable as possible. It would be better to put the code for reading messages and sending messages of these in to their own functions.

7. Finding Nodes

Don't know where to find a node you can connect to? Here are a few places you can try:

- **Your own node.** If you download and run your own Bitcoin Core node on your local computer, you can connect to it at the IP address `127.0.0.1`. Or if you're hosting it on a remote server, use the IP for that server.
- [bitnodes.io](#) This is a handy website that lists of the available nodes on the Bitcoin network that it can find.
- **DNS Seeds.** There are some DNS servers run by trusted Bitcoin Core developers that will return some IPs of reliable full nodes. You can query these DNS seeds by using any online "DNS lookup" tool. Here are some examples of DNS Seeds:
 - `seed.bitcoin.sipa.be` - Pieter Wuille
 - `dnsseed.bitcoin.dashjr.org` - Luke Dashjr
 - `seed.bitcoin.sprovoost.nl` - Sjors Provoost

Tip: You can perform a DNS request on a DNS seed from the command line with: `nslookup seed.bitcoin.sipa.be`. Note that this may not work if you are using a VPN.

Bitcoin Core

In terms of how the *Bitcoin Core* client finds nodes to connect to, it looks for them in the following order when starting up:

1. **Previous Connections.** Bitcoin Core maintains a list of nodes it has previously connected to, and tries connecting to those again once it starts up.
2. **DNS Seeds.** If you're running Bitcoin Core for the first time, you won't have a database of previous nodes, so it will use DNS Seeds like the ones above to find nodes to connect to.
3. **Hardcoded List.** If all else fails, Bitcoin Core comes with a hard-coded list of "Seed Nodes" it will connect to and use as a starting point to help it find other nodes on the network. This list can be found in [chainparamsseeds.h](#).

Ultimately the goal is to just be able to connect to *one* other reliable node on the network, because from there that node will be able to let you know about other nodes you can connect to, and so on and so on.

8. Summary

Connecting to a node from scratch is a cool way to get started with programming in Bitcoin. It allows you to see how nodes communicate with each other, and it gives you live access to the latest transactions and blocks on the network.

You can connect to a node from pretty much any programming language you like. All you need is to be able to make TCP connections and have the IP and port number for a computer running a bitcoin node. If you're running bitcoin locally, the IP will be `127.0.0.1` and the port will be `8333` (by default).

The trickiest part by far is figuring out how to construct messages. You need to get all the raw bytes of data in the correct order, because even if you get one byte wrong, the node you're sending messages will not understand you. And this can be a somewhat frustrating process until you get it right. But once you've got that first message sent correctly, all of the other [message types](#) are much easier to construct.

Getting my first raw transaction from a real-life bitcoin node using a script I wrote from scratch was one of the most satisfying achievements of my programming career.

Good luck.